

ABHIRAM MIRTHIPATI

+91 9440971346 ◊ Mumbai, India
mirthipatiabhiram@gmail.com ◊ linkedin.com/in/abhiram

EDUCATION

Bachelor of Technology, National Institute of Technology, Nagpur 2020 - 2024
CGPA: 8.72 / 10

SKILLS

Programming	Python, SQL, OOPS, C#
Backend & Frameworks	FastAPI, WebSockets, HTTP Protocols, API Performance Optimization
Streaming & Messaging	Apache Kafka, Redis Streams
GenAI & LLMs	LLM Integration, Agentic AI, RAG, Prompt Engineering, Inference Pipelines
Database	MongoDB, Oracle Cloud Database, MySQL, Vector Database, Redis
Cloud	Oracle Cloud Infrastructure (OCI), Google Cloud Platform (GCP)
DevOps & Tools	Git, Linux, Deployment, Postman, Apache JMeter, New Relic Infra, JIRA

EXPERIENCE

Software Engineer (Backend Developer) Oct 2024 – Present
Decimal Point Analytics, India

- Built and scaled Python backend systems handling 35,000+ news articles per day in production.
- Developed high-performance FastAPI REST and WebSocket APIs, reducing latency by 50%.
- Designed real-time data streaming pipelines using Kafka and Redis for low-latency data delivery.
- Integrated LLM-based workflows through OpenAI API calls for automated text processing and content enrichment.
- Built an MCP chatbot server that interprets and validates user queries using NLP, then retrieves relevant data from the database.
- Architected and delivered a distributed security filings conversion platform processing XBRL/iXBRL, HTML, PDF, DOCX, and XLSX filings with real-time API-based delivery via Oracle Object Storage; designed Kafka-driven event pipelines to scale processing across filing types, ensuring high-throughput, low-latency distribution to enterprise customers.

PROJECTS

Hybrid RAG-Based Insurance Claims Correspondence Agent July 2026 – July 2026
Python, Google ADK, LiteLLM, Hybrid RAG

- Built a multi-agent Google ADK and LiteLLM pipeline to draft insurance claim correspondence as .eml files.
- Implemented hybrid BM25 and vector retrieval with reranking.
- Designed guardrail callbacks to block hallucinated claim data and cross-claim information leakage.

PDF Query Application Oct 2025 – Nov 2025
RAG, Vector DB

- Built a PDF Query Application using Retrieval-Augmented Generation (RAG) to enable natural language querying over unstructured documents.
- Implemented text extraction, chunking, and semantic embedding, storing vectors in a database for efficient similarity-based retrieval.

CERTIFICATIONS

Oracle AI Vector Search Certified Professional Sept 2025
Databricks Generative AI Fundamentals Oct 2025